TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

Gaussian Noisy Channel RDAs

A General Autoencoder

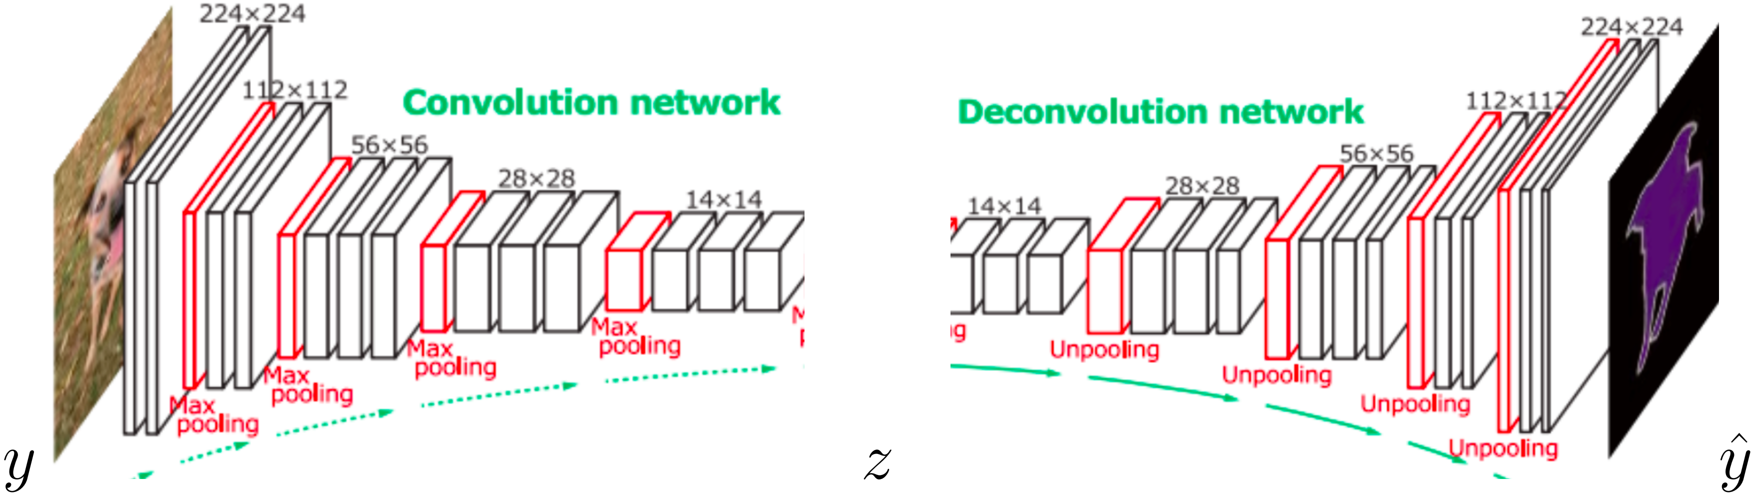

In generall we have either $P_\Phi(z)$ for z discrete or $\hat{p}_\Phi(z)$ for z continuous.

A General Autoencoder

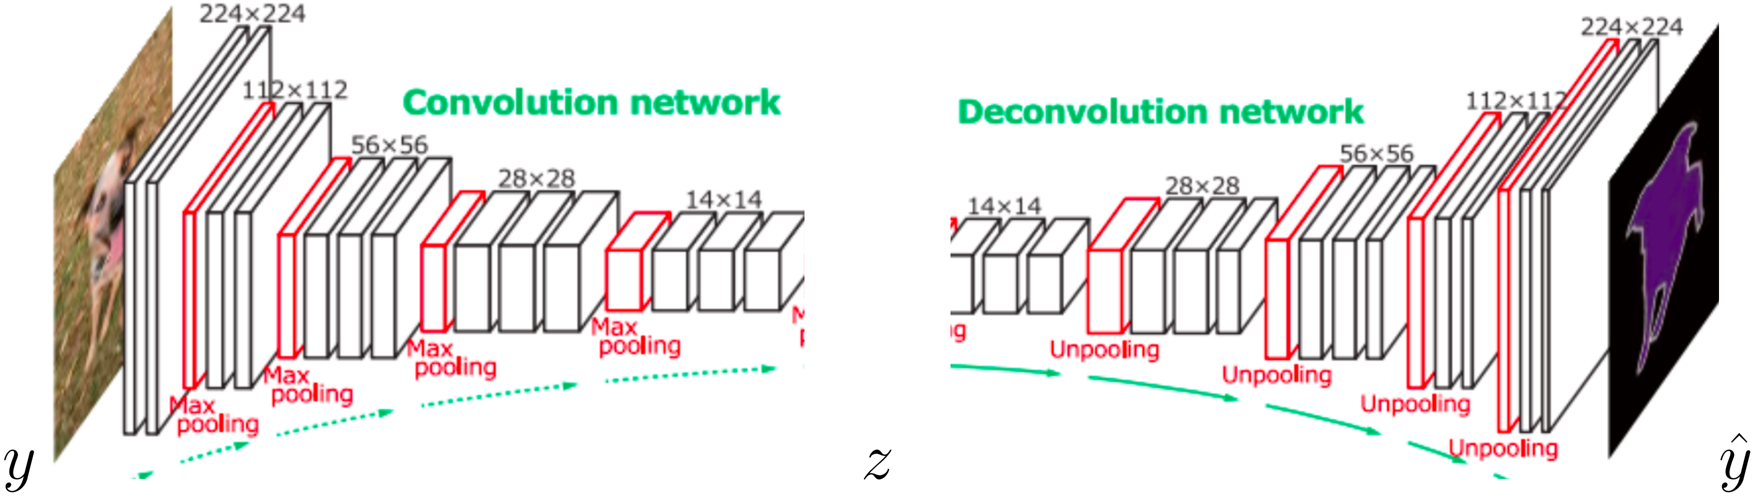

Here we will show that for the continuous case with $p_\Phi(z|y)$ and $\hat{p}_\Phi(z)$ both Gaussian, we can assume without loss of generality that

$$\hat{p}_\Phi(z) = \mathcal{N}(0, I)$$

Gaussian Noisy-Channel RDA

$$\Phi^* = \underset{\Phi}{\text{argmin}}\, E_{y,\epsilon}\ \ln \frac{p_\Phi(z_\Phi(y,\epsilon)|y)}{\hat{p}_\Phi(z_\Phi(y,\epsilon))} + \lambda \text{Dist}(y, y_\Phi(z_\Phi(y,\epsilon)))$$

$$z_\Phi(y,\epsilon) = \mu_\Phi(y) + \sigma_\Phi(y) \odot \epsilon \quad \epsilon \sim \mathcal{N}(0,I)$$

$$p_\Phi(z[i]|y) = \mathcal{N}(\mu_\Phi(y)[i], \sigma_\Phi(y)[i]))$$

$$\hat{p}_\Phi(z[i]) = \mathcal{N}(\hat{\mu}_z[i], \hat{\sigma}_z[i])$$

$$\text{Dist}(y,\hat{y}) = ||y - \hat{y}||^2$$

Gaussian Noisy-Channel RDA

$$\Phi^* = \operatorname*{argmin}_{\Phi} E_{y,\epsilon}\ \ln \frac{p_\Phi(z_\Phi(y,\epsilon)|y)}{\hat{p}_\Phi(z_\Phi(y,\epsilon))} + \lambda \mathrm{Dist}(y, y_\Phi(z_\Phi(y,\epsilon)))$$

We will show that we can fix $\hat{p}_\Phi(z)$ to $\mathcal{N}(0, I)$.

$$p_\Phi(z[i]|y) = \mathcal{N}(\mu_\Phi(y)[i], \sigma_\Phi(y)[i])$$

$$\hat{p}_\Phi(z[i]) = \mathcal{N}(0, 1)$$

$$\mathrm{Dist}(y, \hat{y}) = ||y - \hat{y}||^2$$

Gaussian Noisy-Channel RDA

$$\begin{aligned}\Phi^* &= \operatorname*{argmin}_\Phi E_{y,\epsilon}\ \ln \frac{p_\Phi(z_\Phi(y,\epsilon)|y)}{\hat{p}_\Phi(z_\Phi(y,\epsilon))} + \lambda \mathrm{Dist}(y, y_\Phi(z_\Phi(y,\epsilon))) \\ &= \operatorname*{argmin}_\Phi\ E_{y\sim\mathrm{Pop}} \left(\begin{array}{c} KL(p_\Phi(z|y), \hat{p}_\Phi(z)) \\ \\ +\lambda\ E_\epsilon\ \mathrm{Dist}(y,\ y_\Phi(z_\Phi(y,\epsilon))) \end{array}\right)\end{aligned}$$

Closed Form KL-Divergence

$$KL(p_\Phi(z|y), \hat{p}_\Phi(z))$$

$$= \sum_i \frac{\sigma_\Phi(y)[i]^2 + (\mu_\Phi(y)[i] - \mu_z[i])^2}{2\sigma_z[i]^2} + \ln \frac{\sigma_z[i]}{\sigma_\Phi(y)[i]} - \frac{1}{2}$$

Standardizing $\hat{p}_\Phi(z)$

$$KL(p_\Phi(z|y), p_\Phi(z))$$

$$= \sum_i \frac{\sigma_\Phi(y)[i]^2 + (\mu_\Phi(y)[i] - \mu_z[i])^2}{2\sigma_z[i]^2} + \ln \frac{\sigma_z[i]}{\sigma_\Phi(y)[i]} - \frac{1}{2}$$

$$KL(p_{\Phi'}(z|y), \mathcal{N}(0, I))$$

$$= \sum_i \frac{\sigma_{\Phi'}(y)[i]^2 + \mu_{\Phi'}(y)[i]^2}{2} + \ln \frac{1}{\sigma_{\Phi'}(y)[i]} - \frac{1}{2}$$

Standardizing $\hat{p}_\Phi(z)$

$$KL_\Phi = \sum_i \frac{\sigma_\Phi(y)[i]^2 + (\mu_\Phi(y)[i] - \mu_z[i])^2}{2\sigma_z[i]^2} + \ln \frac{\sigma_z[i]}{\sigma_\Phi(y)[i]} - \frac{1}{2}$$

$$KL_{\Phi'} = \sum_i \frac{\sigma_{\Phi'}(y)[i]^2 + \mu_{\Phi'}(y)[i]^2}{2} + \ln \frac{1}{\sigma_{\Phi'}(y)[i]} - \frac{1}{2}$$

Setting Φ' so that

$$\begin{aligned} \mu_{\Phi'}(y)[i] &= (\mu_\Phi(y)[i] - \mu_z[i])/\sigma_z[i] \\ \sigma_{\Phi'}(y)[i] &= \sigma_\Phi(y)[i]/\sigma_z[i] \end{aligned}$$

gives $KL(p_\Phi(z|y), \hat{p}_\Phi(z)) = KL(p_{\Phi'}(z|y), \mathcal{N}(0, I))$.

Sampling

Sample $z \sim \mathcal{N}(0, I)$ and compute $y_\Phi(z)$

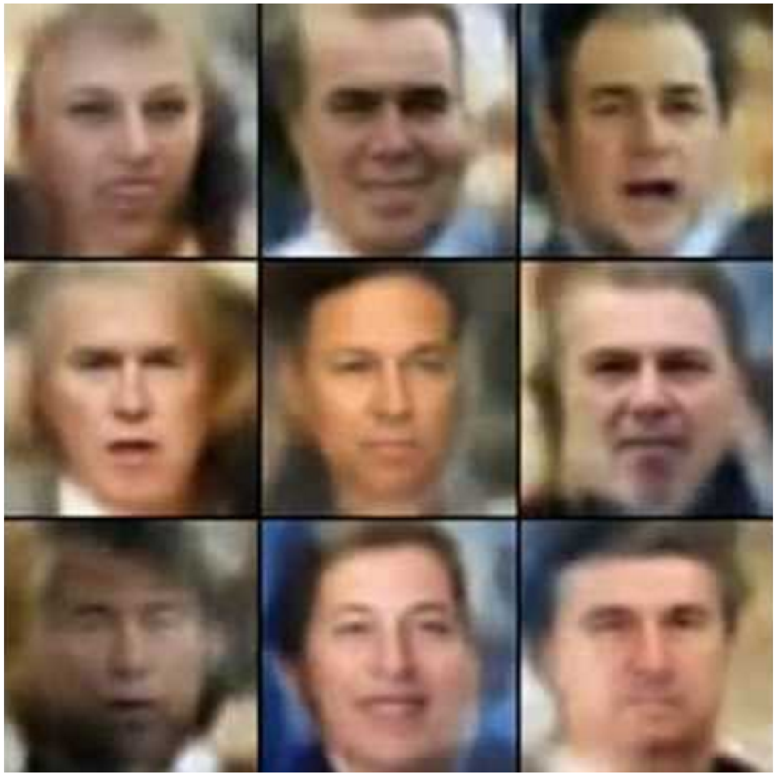

[Alec Radford]

Summary: Gaussian RDAs

Gaussian RDA: $z_\Phi(y,\epsilon) = \mu_\Phi(y) + \sigma_\Phi(y) \odot \epsilon, \qquad \epsilon \sim \mathcal{N}(0,I)$

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}}\, E_{y\sim\mathrm{Pop}} \left(\begin{array}{cc} & KL(p_\Phi(z|y), \mathcal{N}(0,I)) \\ \\ + & \lambda E_\epsilon\, \mathrm{Dist}(y, y_\Phi(z_\Phi(y,\epsilon))) \end{array} \right)$$

END