TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

PAC-Bayesian Learning Theory

Chomsky vs. Kolmogorov and Hinton

Noam Chomsky: Natural language grammar cannot be learned by a universal learning algorithm. This position is supported by the "no free lunch theorem".

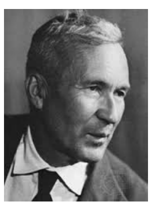

Andrey Kolmogorov, Geoff Hinton: Universal learning algorithms exist. This position is supported by the "free lunch theorem".

The No Free Lunch Theorem

Without prior knowledge, such as universal grammar, it is impossible to make a prediction for an input you have not seen in the training data.

Proof: Select a predictor h uniformly at random from all functions from $\mathcal{X}$ to $\mathcal{Y}$ and then take the data distribution to draw pairs $(x, h(x))$ where x is drawn uniformly from $\mathcal{X}$. No learning algorithm can predict $h(x)$ where x does not occur in the training data.

The Free Lunch Theorem

Consider a classifier f written in C++ with an arbitrarily large standard library.

Let $|f|$ be the number of bits needed to represent f.

The Free Lunch Theorem

$$0 \leq \mathcal{L}(h, x, y) \leq L_{\mathrm{max}}$$

$$\begin{aligned} \mathcal{L}(h) &= E_{(x,y)\sim\mathrm{Pop}}\, \mathcal{L}(h, x, y) \\ \hat{\mathcal{L}}(h) &= E_{(x,y)\sim\mathrm{Train}}\, \mathcal{L}(h, x, y) \end{aligned}$$

Theorem: With probability at least $1 - \delta$ over the draw of the training data the following holds simultaneously for all f.

$$E(f) \leq \frac{10}{9}\left(\hat{E}(f) + \frac{5L_{\mathrm{max}}}{N_{\mathrm{Train}}}\left((\ln 2)|f| + \ln\frac{1}{\delta}\right)\right)$$

Free Lunch Theorem (Probability Form)

Code length is inter-convertable with with probability.

$$P(h) = 2^{-|h|} \quad \text{or} \quad |h| = -\log_2 P(h)$$

Instead of fixing the language (e.g., C++ with a large library) we fix a prior $P(h)$.

Theorem: With probability at least $1 - \delta$ over the draw of training data the following holds simultaneously for all h.

$$\mathcal{L}(h) \leq \frac{10}{9}\left(\hat{\mathcal{L}}(h) + \frac{5L_{\max}}{N_{\mathrm{Train}}}\left(\ln\frac{1}{P(h)} + \ln\frac{1}{\delta}\right)\right)$$

Proof

Define

$$\epsilon(h) = \sqrt{\frac{2\mathcal{L}(h)\left(\ln\frac{1}{P(h)} + \ln\frac{1}{\delta}\right)}{L_{\max}N_{\mathrm{Train}}}}.$$

By the relative Chernov bound we have

$$P_{\mathrm{Train}\sim\mathrm{Pop}}\left(\frac{\hat{\mathcal{L}}(h)}{L_{\max}} \leq \frac{\mathcal{L}(h)}{L_{\max}} - \epsilon(h)\right) \leq e^{-N_{\mathrm{Train}}\frac{\epsilon(h)^2 L_{\max}}{2\mathcal{L}(h)}} = \delta P(h).$$

Proof

$$P_{\mathrm{Train}\sim\mathrm{Pop}}\left(\hat{\mathcal{L}}(h) \leq \mathcal{L}(h) - L_{\max}\epsilon(h)\right) \leq \delta P(h).$$

$$P_{\mathrm{Train}\sim\mathrm{Pop}}\left(\exists h\ \hat{\mathcal{L}}(h) \leq \mathcal{L}(h) - L_{\max}\epsilon(h)\right) \leq \sum_h \delta P(h) = \delta$$

$$P_{\mathrm{Train}\sim\mathrm{Pop}}\left(\forall h\ \mathcal{L}(h) \leq \hat{\mathcal{L}}(h) + L_{\max}\epsilon(h)\right) \geq 1 - \delta$$

Proof

$$\mathcal{L}(h) \leq \widehat{\mathcal{L}}(h) + L_{\max}\sqrt{\frac{\mathcal{L}(h)}{L_{\max}}\left(\frac{2\left(\ln\frac{1}{P(h)} + \ln\frac{1}{\delta}\right)}{N_{\mathrm{Train}}}\right)}$$

using

$$\sqrt{ab} = \inf_{\lambda>0}\ \frac{a}{2\lambda} + \frac{\lambda b}{2}$$

we get

$$\mathcal{L}(h) \leq \widehat{\mathcal{L}}(h) + \frac{\mathcal{L}(h)}{2\lambda} + \frac{\lambda L_{\max}\left(\ln\frac{1}{P(h)} + \ln\frac{1}{\delta}\right)}{N_{\mathrm{Train}}}$$

Proof

$$\mathcal{L}(h) \leq \widehat{\mathcal{L}}(h) + \frac{\mathcal{L}(h)}{2\lambda} + \frac{\lambda L_{\max}\left(\ln \frac{1}{P(h)} + \ln \frac{1}{\delta}\right)}{N_{\text{Train}}}$$

Solving for $\mathcal{L}(h)$ yields

$$\mathcal{L}(h) \leq \frac{1}{1 - \frac{1}{2\lambda}}\left(\hat{\mathcal{L}}(h) + \frac{\lambda L_{\max}}{N_{\text{Train}}}\left(\ln \frac{1}{P(h)} + \ln \frac{1}{\delta}\right)\right)$$

Setting $\lambda = 5$ brings the leading factor to $10/9$ which seems sufficiently close to 1 that larger values of λ need not be considered.

A Model Compression Guarantee

Let $|\Phi|$ be the number of bits used to represent Φ under some fixed compression scheme.

Let $P(\Phi) = 2^{-|\Phi|}$

$$\mathcal{L}(\Phi) \leq \frac{10}{9}\left(\hat{\mathcal{L}}(\Phi) + \frac{5L_{\mathrm{max}}}{N_{\mathrm{Train}}}\left((\ln 2)|\Phi| + \ln\frac{1}{\delta}\right)\right)$$

A Bound for Continuous Densities

Let p be any "prior" and q be any "posterior" on any (possibly continuous) model space. Define

$$L(q) = E_{h \sim q}\, L(h)$$

$$\hat{L}(q) = E_{h \sim q}\, \hat{L}(h)$$

For any p and any $\lambda > \frac{1}{2}$, with probability at least $1-\delta$ over the draw of the training data, the following holds simultaneously for all q.

$$L(q) \leq \frac{1}{1 - \frac{1}{2\lambda}} \left(\hat{L}(q) + \frac{\lambda L_{\max}}{N_{\mathrm{Train}}} \left(KL(q,p) + \ln \frac{1}{\delta} \right) \right)$$

Adding Noise Simulates Limiting Precision

Assume $0 \leq \mathcal{L}(\Phi, x, y) \leq L_{\max}$.

Define:

$$\mathcal{L}(\Phi) = E_{(x,y)\sim\text{Pop},\ \epsilon\sim\mathcal{N}(0,\sigma)^d}\ \mathcal{L}(\Phi + \epsilon, x, y)$$

$$\hat{\mathcal{L}}(\Phi) = E_{(x,y)\sim\text{Train},\ \epsilon\sim\mathcal{N}(0,\sigma)^d}\ \mathcal{L}(\Phi + \epsilon, x, y)$$

Theorem: With probability at least $1 - \delta$ over the draw of training data the following holds **simultaneously** for all Φ.

$$\mathcal{L}(\Phi) \leq \frac{10}{9}\left(\hat{\mathcal{L}}(\Phi) + \frac{5L_{\max}}{N_{\text{Train}}}\left(\frac{||\Phi - \Phi_{\text{init}}||^2}{2\sigma^2} + \ln\frac{1}{\delta}\right)\right)$$

Non-Vacuous Generalization Guarantees

Model compression has recently been used to achieve "non-vacuous" PAC-Bayes generalization guarantees for ImageNet classification — error rate guarantees less than 1.

Non-Vacuous PAC-Bayes Bounds at ImageNet Scale.

Wenda Zhou, Victor Veitch, Morgane Austern, Ryan P. Adams, Peter Orbanz

ICLR 2019

END